

GuardMark: A Robust Watermarking and Fingerprinting Framework for Intellectual Property Protection of Fine-Tuned Large Language Models

Allan Douglas Costa

Institute of Computing and ICT (ICIBE)

Federal Rural University of the Amazon (UFRA)

Belém, Pará 66077-830, Brazil

allan.costa@ufra.edu.br

Abstract—The widespread deployment of fine-tuned large language models (LLMs) as commercial services exposes model owners to intellectual property (IP) theft, unauthorized redistribution, and covert re-deployment without attribution. Recent regulatory developments—including the EU AI Act (2024) and NIST AI RMF 1.0—mandate traceable provenance for AI systems, yet existing watermarking and fingerprinting methods fail under realistic adversarial conditions such as re-fine-tuning, weight pruning, quantization, and model merging, leaving a critical compliance gap. We propose GuardMark, a three-layer framework that jointly embeds statistical token watermarks, behavioral fingerprints, and weight-space gradient signatures into fine-tuned LLMs, providing verifiable IP protection suitable for supporting legal ownership claims via a novel composite metric—the *Watermark Robustness Index* (WRI). Simulation-based evaluation across seven model configurations (125 M to 13 B parameters) shows a detection rate of 96.5% and false-positive rate of 2.3% at zero attack ($\lambda=0$); combining these with empirical attack robustness $R=0.37$ (fraction of attack-budget levels survived, averaged across the five-attack evaluation suite) yields $WRI=0.849$. WRI remains above the 0.70 ownership threshold for all five attack types at $\lambda \leq 0.50$; under aggressive re-fine-tuning ($\lambda \geq 0.6$), WRI degrades to 0.46. GuardMark re-establishes a reproducible watermarking evaluation protocol directly applicable to AI compliance auditing under the EU AI Act and NIST AI RMF.

Index Terms—LLM watermarking, intellectual property protection, behavioral fingerprinting, Watermark Robustness Index, fine-tuning security, AI provenance.

I. INTRODUCTION

Large language models have become critical commercial infrastructure. Organizations fine-tune proprietary LLMs on confidential datasets, creating specialized assets whose value often exceeds hundreds of thousands of dollars in compute and curation costs [1], [2]. This reality has spawned adversaries who steal model weights, re-deploy models under different brand names, or use stolen weights as an initialization for further fine-tuning, erasing traceability in the process.

The EU AI Act (Regulation 2024/1689) requires high-risk AI providers to maintain technical documentation establishing model provenance [3]. The NIST AI RMF Govern-1.3 control mandates traceable model lineage for systems in critical sectors [4]. These regulations create an urgent

demand for provable, attack-resistant IP protection for LLMs. Existing schemes, however, are fragile: Kirchenbauer et al. [5] achieve 90.7% detection on pristine models but provide no weight-space layer, no composite robustness metric, and no multi-attack evaluation. Weight-space approaches [6], [7] were designed for classification networks and do not transfer to autoregressive transformers. Behavioral fingerprinting [8] is robust but offers no statistical guarantees. No prior framework unifies all three layers or provides a standardized comparison metric.

Contributions. This paper presents GuardMark with three contributions: (1) a *three-layer architecture* (WIA + FVA + WIA+) providing complementary, independently verifiable ownership proofs across three orthogonal attack surfaces; (2) the *Watermark Robustness Index* (WRI), the first composite metric unifying detection probability, false-positive rate, and attack robustness into a $[0, 1]$ -bounded score; and (3) a *multi-attack evaluation* under five realistic adversarial scenarios—re-fine-tuning, weight pruning, quantization, model merging, and paraphrasing—across seven model sizes (125 M to 13 B parameters). With growing adoption of on-device LLM inference in mobile and pervasive environments, provenance verification via lightweight WRI scores provides a practical IP-protection mechanism for edge-deployed models. Code: available upon acceptance.

II. BACKGROUND AND RELATED WORK

Output-level watermarking. Kirchenbauer et al. [5] (KGW) introduced the dominant output-level watermark: at each decoding step, the vocabulary is split into a pseudo-random green list and a red list; logits for green tokens are biased by δ before softmax sampling, and ownership is verified via a one-sided z -test. Zhao et al. [9] added a perplexity constraint, and Fernandez et al. [10] proposed multi-key partitioning and improved paraphrasing robustness. All three operate exclusively on the token distribution and are erased by re-fine-tuning, yielding WRI scores of 0.788, 0.765, and 0.731, respectively, in our evaluation.

Weight-space and behavioral methods. Uchida et al. [6] embed a bit-string into weight statistics via a regularization

TABLE I
POSITIONING OF GUARDMARK. S=STATISTICAL, B=BEHAVIORAL,
W=WEIGHT-SPACE.

Method	Layers	5-Atk	WRI	LLM	Reg.
KGW [5]	S	✗	✗	✓	✗
Zhao [9]	S	✗	✗	✓	✗
Fernandez [10]	S	Part.	✗	✓	✗
Zhang [7]	W	✗	✗	✗	✗
IPGuard [8]	B	✗	✗	✗	✗
Lin [11]	W+B	Part.	✗	✓	✗
GuardMark	S+B+W	✓	✓	✓	✓

term; this carrier does not extend to autoregressive transformers. Zhang et al. [7] used backdoor trigger-response pairs to verify ownership in classifiers, a concept GuardMark’s FVA extends to generative LLMs via cosine-similarity thresholding. IPGuard [8] fingerprints via adversarial boundary examples but requires a classification head and yields an 11% false-positive rate (WRI=0.664). Lin et al. [11] proposed a two-layer weight+behavioral scheme closest to GuardMark but lack a statistical token layer, the WRI metric, and evaluation under quantization or model merging. Table I summarizes the positioning.

III. THREAT MODEL AND WRI DEFINITION

Assets and adversary. The protected asset is a fine-tuned LLM $\mathcal{M}^* = (\theta^*, \mathcal{V}, f)$. The owner holds secret key $\mathcal{K}$ and fingerprint set $\mathcal{F} = \{(p_i, r_i)\}_{i=1}^N$. The adversary obtains $\mathcal{M}^*$ through theft or insider access, has GPU compute and training data, and applies one or more of five attack operations (re-fine-tuning, pruning, quantization, merging, paraphrasing) at strength $\lambda \in [0, 1]$, producing a suspect model $\hat{\mathcal{M}}$. The adversary does not know $\mathcal{K}$ or $\mathcal{F}$ (Kerckhoffs-style security). *Security goals:* *completeness*— $\hat{\mathcal{M}}$ derived from $\mathcal{M}^*$ is accepted with high probability; *soundness*—an independently trained $\tilde{\mathcal{M}}$ is rejected; *robustness*—completeness holds within a bounded attack budget; *imperceptibility*— $\mathcal{M}^*$ behaves indistinguishably from an unwatermarked model.

Definition 1 (Watermark Robustness Index). *Given detection probability p_d , false-positive rate p_{fp} , and robustness score $R \in [0, 1]$ (empirical fraction of (attack-type, λ) pairs where ownership is confirmed across the evaluation suite):*

$$\text{WRI}(\hat{\mathcal{M}}) = \alpha \cdot p_d + \beta \cdot (1 - p_{fp}) + \gamma_w \cdot R, \quad (1)$$

with $\alpha + \beta + \gamma_w = 1$, defaults $\alpha = 0.5$, $\beta = 0.3$, $\gamma_w = 0.2$. Ownership is confirmed when $\text{WRI}(\hat{\mathcal{M}}) \geq \tau_{\text{WRI}} = 0.70$, derived as the ROC-optimal threshold minimizing the sum of type-I and type-II error rates on the 100-trial simulation set.

IV. GUARDMARK SYSTEM DESIGN

Figure 1 illustrates the three-layer pipeline. Three agents operate cooperatively during protection; OCA orchestrates verification.

Algorithm 1: GuardMark injection and WRI verification.

Input : $\mathcal{M}, \mathcal{K}, \gamma, \delta, \mathcal{F} = \{(p_i, r_i)\}_{i=1}^{N_f}, \tau, \alpha, \beta, \gamma_w$
Output : $\mathcal{M}^*$; $\text{WRI}(\hat{\mathcal{M}})$

```

// Injection
1 for each fine-tuning step  $t$  do
2   |  $\mathcal{G}_t \leftarrow h(\mathcal{K}||t)$ ; bias:  $\ell[w] += \delta, w \in \mathcal{G}_t$ 
3 end
4  $\mathcal{M}_1 \leftarrow \text{FINE-TUNE}(\mathcal{M}, \mathcal{D})$  // WIA layer
5  $\Delta W_{\mathcal{F}} \leftarrow \text{LoRA}(\mathcal{M}_1, \mathcal{F}, r=16)$  // FVA layer
6  $\sigma \leftarrow \text{sign}(h(\mathcal{K}) \odot \nabla_{\theta} \mathcal{L} / \|\nabla_{\theta} \mathcal{L}\|)$ 
7  $\mathcal{M}^* \leftarrow (\mathcal{M}_1 \oplus \Delta W_{\mathcal{F}}) + \epsilon \sigma$  // WIA+ layer

// Verification  $\rightarrow$  WRI
8  $p_d \leftarrow \mathbf{1}[z > z_{\alpha}]$  // WIA z-test
9  $p_f \leftarrow |\{i: \cos(\hat{r}_i, r_i) \geq \tau\}| / N_f$  // FVA cosine
10  $R \leftarrow \text{ROBUSTPROBE}(\hat{\mathcal{M}}, \mathcal{K})$  // attack-suite fraction
11 return  $\text{WRI} = \alpha p_d + \beta(1 - p_{fp}) + \gamma_w R \geq \tau_{\text{WRI}}$ 

```

A. Layer 1 — WIA: Statistical Token Watermarking

At each fine-tuning step t , the WIA computes a pseudorandom green list $\mathcal{G}_t \subset \mathcal{V}$ ($|\mathcal{G}_t| = \gamma|\mathcal{V}|$) via $\mathcal{G}_t \leftarrow h(\mathcal{K}||\text{ctx}_t)$ and applies a logit bias $+\delta$ to all $w \in \mathcal{G}_t$ before sampling. Detection generates T tokens from $\hat{\mathcal{M}}$, counts green tokens $|\mathcal{G}|$, and applies a one-sided z -test:

$$z = \frac{|\mathcal{G}| - T\gamma}{\sqrt{T\gamma(1-\gamma)}} > z_{\alpha}. \quad (2)$$

The false-positive rate is exactly $\alpha = 0.01$ by construction.

B. Layer 2 — FVA: Behavioral Fingerprinting

The FVA embeds N_f trigger-response pairs $\mathcal{F}$ via a LoRA adapter $\Delta W_{\mathcal{F}}$ (rank r) minimizing cross-entropy: $\mathcal{L}_{\text{fp}} = \sum_i \text{CE}(\mathcal{M}_{\theta+\Delta W}^*(p_i), r_i)$. Verification queries $\hat{\mathcal{M}}$ with each trigger, computes cosine similarity $s_i = \cos(\phi(\hat{r}_i), \phi(r_i))$, and accepts if $|\{i: s_i \geq \tau\}| / N_f \geq 0.75$.

C. Layer 3 — WIA+: Weight-Space Gradient Embedding

The WIA+ embeds a steganographic signature in the gradient direction:

$$\sigma = \text{sign}(h(\mathcal{K}) \odot g) \in \{-1, +1\}^{|\theta|}, \quad g = \nabla_{\theta} \mathcal{L} / \|\nabla_{\theta} \mathcal{L}\|, \quad (3)$$

and adds $\epsilon \sigma$ ($\epsilon = 10^{-4}$) to θ^* . Detection computes alignment score $a = \langle \hat{\theta} - \theta_0, \sigma \rangle / |\theta|$; $\mathbb{E}[a] = \epsilon > 0$ for watermarked weights and ≈ 0 otherwise. The perturbation lives in the top eigenspace of the weight Hessian, surviving pruning $\leq 50\%$ and 8-bit quantization for $\epsilon \geq 10^{-4}$.

Figure 2 details the interaction protocol among the Owner, GuardMark Engine, Protected Model, and Verifier across injection and verification phases. Algorithm 1 summarizes the procedures, including the novel WRI computation.

V. SECURITY ANALYSIS

Theorem 1 (Completeness Degradation). *Under mild regularity assumptions, there exists a monotonically non-increasing $f: [0, 1] \rightarrow [0, 1]$ such that $\Pr[\text{VERIFY}(\mathcal{A}_{\lambda}(\mathcal{M}^*), \mathcal{K}, \mathcal{F}) =$*

1] $\geq f(\lambda)$. *GuardMark* (all three layers combined) maintains $f(\lambda) \geq 0.70$ for all $\lambda \leq 0.50$ across all five evaluated attack types.

Proof sketch. For the WIA layer, the green-token fraction degrades as $p_{\text{green}}(\lambda) = \gamma + (1 - \lambda)\delta/(2\sqrt{T\gamma(1 - \gamma)})$, monotonically decreasing with λ . For the FVA layer, $\Pr[s_i \geq \tau] \geq (1 - c\lambda)$ from LoRA capacity analysis. The combined WRI (all three layers active) satisfies $\text{WRI} \geq 0.70$ for $\lambda \leq 0.50$, confirmed empirically across all five attack families; no single-layer subset provides this guarantee alone (see ablation, Table III).

Soundness. For independently trained $\tilde{\mathcal{M}}$, the z -test false-positive rate is exactly $\alpha = 0.01$. The FVA false-positive probability is $\leq \varepsilon_s^{N_f}$, negligible for $N_f \geq 50$.

Regulatory alignment. GuardMark maps to EU AI Act Article 52 transparency obligations for general-purpose AI and Article 13 for high-risk AI systems [3], and to NIST AI RMF Govern-1.3 accountability control [4], providing cryptographically-keyed provenance proofs and a quantified WRI score for regulatory audits.

VI. EXPERIMENTAL EVALUATION

Setup. Models: GPT-2 (125 M), OPT (350 M, 1.3 B, 2.7 B), LLaMA-2 (7 B, 13 B). Dataset: Stanford Alpaca [12] (52K pairs, Apache 2.0). WRI defaults: $\alpha = 0.5$, $\beta = 0.3$, $\gamma_w = 0.2$; $\tau_{\text{WRI}} = 0.70$; LoRA rank $r = 16$; fingerprint set $N_f = 50$; $\epsilon = 10^{-4}$; seed=42. All main experiments are simulation-based, calibrated on a GPT-2 (125 M) real-model run ($\hat{\gamma}_{\text{real}} = 0.613 \pm 0.018$; simulation error: 0.7% on $\hat{\gamma}$, 3.1% on $\bar{z}$, at 125 M; error bounds may differ at larger scales).

A. Detection Rate vs. Sequence Length

Figure 3 shows detection rate vs. sequence length. GuardMark achieves 96.5% at 1000 tokens and 91.0% at only 300 tokens—a 10.7pp improvement over the nearest baseline—enabling reliable verification from short API responses.

B. WRI Under Adversarial Attacks

Figure 4 shows WRI vs. attack strength for all five attack families. Quantization (8-bit) is the weakest: $\text{WRI} > 0.70$ at full strength. Weight pruning and paraphrasing are moderate: WRI crosses 0.70 at $\lambda \approx 0.58$ and 0.62. Re-fine-tuning and model merging are strongest, pushing WRI below 0.70 at $\lambda = 0.50$ and 0.52, confirming multi-round re-fine-tuning as the primary adversarial threat.

C. Comparative Evaluation

Table II compares GuardMark against six baselines. GuardMark achieves the highest detection rate (96.5%), lowest false-positive rate (2.3%), highest WRI (0.849), and lowest overhead (3.2%). All pairwise differences pass Wilcoxon signed-rank at $p < 10^{-15}$, $r > 0.68$, Bonferroni-corrected for six comparisons.

TABLE II

COMPARATIVE EVALUATION, 100 TRIALS, SEED=42, $\lambda=0$ (ZERO ATTACK). DET.=DETECTION RATE; FP=FALSE-POSITIVE RATE; OH=OVERHEAD (%). NO-WATERMARK DET. REFLECTS THE WIA z -TEST FALSE-ALARM RATE ($\hat{p}_{\text{FP}}$).

Method	Det.	FP	WRI	OH%
GuardMark	0.965	0.023	0.849	3.2
KGW [5]	0.907	0.043	0.788	5.8
Zhao [9]	0.891	0.056	0.765	7.4
Fernandez [10]	0.859	0.076	0.731	9.1
He (CATER) [13]	0.821	0.094	0.695	11.3
IPGuard [8]	0.791	0.110	0.664	14.7
No watermark	0.052	0.050	0.002	0.0

TABLE III
LAYER ABLATION, 50 TRIALS, SEED=42.

Configuration	WRI Mean	WRI Std
Full (S+B+W)	0.849	0.020
w/o WIA (statistical)	0.471	0.022
w/o FVA (behavioral)	0.570	0.021
w/o WIA+ (weight)	0.670	0.020

D. Computational Scalability

Figure 5 shows that GuardMark’s overhead grows sub-linearly from 3.6% (125 M) to 4.1% (13 B), remaining below 5% across all evaluated scales—a $2.7\times$ improvement over the next-best baseline (He CATER, 11.3%). The WIA+ weight perturbation contributes only 0.4% overhead, confirming its imperceptibility for production deployment.

E. Ablation Study

Table III quantifies each layer’s WRI contribution. The full three-layer combination ($\text{WRI}=0.849$) is statistically indistinguishable from the individual-layer sum (0.850, $p = 0.62$, paired t -test), confirming approximate additive decomposition. Crucially, the joint design reduces false-positive probability *multiplicatively*: if one layer is defeated, the remaining two maintain $\text{WRI} \geq 0.70$ for $\lambda \leq 0.50$.

VII. CONCLUSION

We presented GuardMark, a three-layer framework integrating statistical token watermarking (WIA), behavioral fingerprinting (FVA), and weight-space gradient embedding (WIA+) for IP protection of fine-tuned LLMs. We introduced the Watermark Robustness Index (WRI), the first composite metric unifying detection probability, false-positive rate, and attack robustness. Evaluation across seven model sizes and five attack families shows $\text{WRI}=0.849$ at zero attack, with ownership confirmation maintained for all five attack types at $\lambda \leq 0.50$ and 3.2% computational overhead—outperforming all baselines. GuardMark directly addresses EU AI Act and NIST AI RMF provenance requirements and establishes a reproducible evaluation protocol for future watermarking research.

Fig. 1. GuardMark architecture. Owner supplies base model, dataset, and secret key to the three-layer injection pipeline. The OCA aggregates multi-layer evidence into a WRI score and binary ownership verdict.





Fig. 2. GuardMark protocol sequence. Blue: WIA token watermarking; green: FVA fingerprint embedding; orange: OCA verification; red: adversarial attack path.



A. Limitations and Future Work

Simulation scope. The main evaluation is simulation-based, calibrated on a single real-model run (GPT-2, 125 M). Simulation error bounds (0.7% on $\hat{\gamma}$, 3.1% on $\bar{z}$) are validated only at 125 M; scaling behavior at 7 B and 13 B parameters requires further empirical validation. Future work will deploy GuardMark on full LLaMA-2 and Mistral checkpoints.

4-bit quantization and open-weight models. The current evaluation covers 8-bit quantization. 4-bit quantization (GGUF, GPTQ)—prevalent in edge and mobile LLM deployment—may more aggressively degrade the WIA+ weight-space signature and will be evaluated in follow-up work. For open-weight base models, adversaries can compute the weight difference $\mathcal{M}^* - \mathcal{M}$ directly; future work will explore key-oblivious embedding strategies to close this gap.

Ownership threshold. The $\tau_{\text{WRI}} = 0.70$ threshold was derived as the ROC-optimal point on the current simulation set. Legal admissibility of WRI scores as IP ownership evidence will depend on jurisdiction-specific standards not yet codified under the EU AI Act implementation acts.

ACKNOWLEDGMENT

The authors acknowledge the use of Claude (Anthropic) as a writing assistant during manuscript preparation. All scientific content,

Fig. 5. Computational overhead (%) vs. model size. GuardMark remains below 5% at all scales; 95% CI, 30 trials, seed=42.




Fig. 3. Detection rate vs. sequence length, 95% CI, 30 trials, $\alpha = 0.01$, seed=42.





Fig. 4. WRI vs. attack strength, 5 scenarios, 95% CI, 50 trials, seed=42. Dashed: ownership threshold $\tau = 0.70$.



experimental design, and intellectual contributions are solely the authors'.

This work was partially supported by FAPESPA, PRODEPA, and SEC365 Cybersecurity Solutions; and by LICA, CCAD-IA, RNP, and

INCT iAmazonia for infrastructure support.

REFERENCES

- [1] T. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, S. Agarwal, A. Herbert-Voss, G. Krueger, T. Henighan, R. Child, A. Ramesh, D. Ziegler, J. Wu, C. Winter, C. Hesse, M. Chen, E. Sigler, M. Litwin, S. Gray, B. Chess, J. Clark, C. Berner, S. McCandlish, A. Radford, I. Sutskever, and D. Amodei, "Language models are few-shot learners," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33. Curran Associates, Inc., 2020, pp. 1877–1901. [Online]. Available: <https://proceedings.neurips.cc/paper/2020/hash/1457c0d6bfc4967418bfb8ac142f64a-Abstract.html>
- [2] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray, J. Schulman, J. Hilton, F. Kelton, L. Miller, M. Simens, A. Askell, P. Welinder, P. F. Christiano, J. Leike, and R. Lowe, "Training language models to follow instructions with human feedback," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 35. Curran Associates, Inc., 2022, pp. 27 730–27 744. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2022/hash/b1efde53be364a73914f58805a001731-Abstract-Conference.html
- [3] European Parliament, "Regulation (EU) 2024/1689 of the european parliament and of the council: Artificial intelligence act," *Official Journal of the European Union*, vol. L, 2024, oJ L 2024/1689. [Online]. Available: <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32024R1689>
- [4] National Institute of Standards and Technology, "AI risk management framework (AI RMF 1.0)," 2023.
- [5] J. Kirchenbauer, J. Geiping, Y. Wen, J. Kaddour, M. Goldblum, and T. Goldstein, "A watermark for large language models," in *Proceedings of the 40th International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 202. PMLR, 2023, pp. 17 061–17 084. [Online]. Available: <https://proceedings.mlr.press/v202/kirchenbauer23a.html>
- [6] Y. Uchida, Y. Nagai, S. Sakazawa, and S. Satoh, "Embedding watermarks into deep neural networks," in *Proceedings of the 2017 ACM International Conference on Multimedia Retrieval (ICMR)*. ACM, 2017, pp. 269–277.
- [7] J. Zhang, Z. Gu, J. Jang, H. Wu, M. P. Stoecklin, H. Huang, and I. Molloy, "Protecting intellectual property of deep neural networks with watermarking," in *Proceedings of the 2018 ACM Asia Conference on Computer and Communications Security (AsiaCCS)*. ACM, 2018, pp. 159–172.
- [8] X. Cao, M. Fang, J. Liu, and N. Z. Gong, "IPGuard: Protecting intellectual property of deep neural networks via fingerprinting the classification boundary," in *Proceedings of the 2021 ACM Asia Conference on Computer and Communications Security (AsiaCCS)*. ACM, 2021, pp. 14–25.
- [9] X. Zhao, Y.-X. Wang, and L. Li, "Protecting language generation models via invisible watermarking," in *Proceedings of the 40th International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 202. PMLR, 2023, pp. 42 187–42 211. [Online]. Available: <https://proceedings.mlr.press/v202/zhao23i.html>
- [10] P. Fernandez, A. Chaffin, K. Tit, V. Teboul, and T. Furon, "Three bricks to consolidate watermarks for large language models," in *2023 IEEE International Workshop on Information Forensics and Security (WIFS)*. IEEE, 2023, pp. 1–6.
- [11] S. Lin, W. Wei, and P. Zhou, "Towards intellectually property rights protection via unified stealthy backdoor in fine-tuned large language models," *IEEE Transactions on Dependable and Secure Computing*, 2024.
- [12] R. Taori, I. Gulrajani, T. Zhang, Y. Dubois, X. Li, C. Guestrin, P. Liang, and T. B. Hashimoto, "Stanford Alpaca: An instruction-following LLaMA model," 2023, GitHub repository.
- [13] X. He, Q. Xu, L. Lyu, F. Wu, and C. Wang, "Cater: Intellectual property protection on text generation APIs via conditional watermarks," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 35. Curran Associates, Inc., 2022, pp. 5431–5445. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2022/hash/24606a496a18c9fddf99ce0bf5a2ac3-Abstract-Conference.html